

Predicting the Quality of Wine

Ethan Evans and Calvin Spencer

{etevans, caspencer}@davidson.edu

Davidson College

Davidson, NC 28035

U.S.A.

Abstract

Wine preferences vary between consumers because of inherent differences in the human brain and palate. In this research, we aim to generate a regression model that can predict how a consumer will grade a wine based on physicochemical characteristics of the wine. While some of these properties have more importance to consumers' ratings than others, it is difficult to derive a model that can accurately predict wine ratings using multivariate analysis. Despite this difficulty, there is a distinct difference between the prediction accuracy on red wines versus white wines. It is clear that more research and experimentation should be done before deploying an algorithm like this to the industrial world, with which its impacts will be commensurate with its accuracy.

1 Introduction

Wine tasting has been a tradition of wine enthusiasts for centuries. Each taster evaluates the quality of the wine based on their own unique palate and generates consumer preferences. The process of rating wines is thought to be completely subjective, but what if one could predict how a particular wine is rated based on physicochemical characteristics of the wine like alcohol content and acidity. Employing this technique would revolutionize the wine consumer market through the creation of new wines and targeted advertising of certain wines.

An idea like that of predicting wine quality based on its characteristics becomes feasible when taking a regression modeling approach. Regression analysis will consider both the characteristics and ratings of pre-existing wines and use patterns in that data to rate wines that the model has not seen yet. Through this type of modeling, we can ideally rate any kind of wine both in existence and in imagination. In this paper, we followed this regression approach to see if we could accurately predict a subjective wine rating from objective inherent characteristics and previous ratings.

The data used to fuel our regression model came from the UCI Machine Learning Repository, created by Paulo Cortez, Antonio Cerdeira, Fernando Almeida, Telmo Matos, and Jose Reis. Data was split into two sets, one set contained data on the Portuguese "Vinho Verde" red wine variant and the other on the corresponding white wine

variant. The following are the eleven physicochemical variables that were recorded for each wine: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, and alcohol. In addition, wine experts graded each wine on a scale of 0 to 10 with 0 representing "very bad" and 10 representing "very excellent." The median of three experts' grades was compiled as this variable. The red wine dataset contains 1,599 wines, and the white wine dataset contains 4,898 wines.

In this paper, we will detail the process of exploring the data, augmenting the data, generating and optimizing regression models for both red and white wines, and comparing the models' effectiveness.

2 Data Preparation

The following data preparation and preprocessing steps were completed independently and identically in both the red wine and white wine data sets. All experimental methodologies and data manipulation was completed using auxiliary Python libraries including *scikit-learn*, *numpy*, *pandas*, and *matplotlib*. The data was first shuffled, and one-fifth of it was held out for a final testing set. The remaining 80% of the data was used for training our regression models.

We plotted individual features against the wine quality to visualize if there were any apparent relationships between them (seen below). As we expected, this varied depending on which feature was chosen. Some features in the red wine dataset, such as alcohol and volatile acidity, showed a directional preference as wine quality increased. This indicated that indeed, some features appear to correlate with a higher quality wine. For others, such as residual sugar, there was no obvious correlation between the feature and the wine quality.

This exploratory phase of our experiment was crucial to the implementation of our solution. These findings indicated some features might be more correlated than others, and led us to consider using a specific type of regression which focuses on reducing the impact of unrelated variables

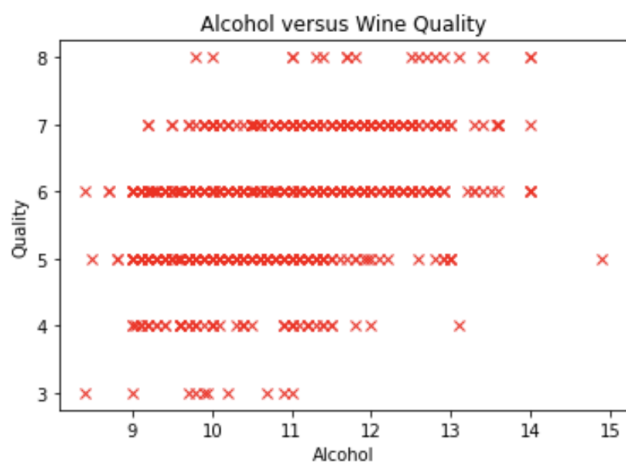


Figure 1: This distribution of wine qualities and their alcohol content indicates a slight positive correlation between the two variables.

on our model.

The eleven input features of the dataset were then scaled using standardization. This transforms the data into a normal distribution feature-wise, which will allow the regression solver to more efficiently converge on a solution. More on this can be found in Scikit-Learn’s data preprocessing user guide (Buitinck et al. 2013).

With all of this background information, we made several key choices with our regression models that allow for the more efficient and accurate prediction model.

3 Experiments

We approached this regression problem knowing that we wanted to employ regularization methods in some manner. Regularization simplifies the model by tuning the weights of certain features, and thus reducing the potential for overfitting. Our initial thought was to implement lasso regression, which is particularly successful in multivariate models where some features may play a larger role in predicting the outcome than others. If you visualize each feature holding some weight or coefficient value, lasso regression shrinks these weights all the way to zero on less important features through a process called L1 regularization. Therefore, due to our prior data exploration process, we hypothesized that some variables would have stronger correlations with wine rating than others, so feature selection was an important consideration for our choice of model.

We compared this to another common regression technique known as ridge regression which implements a slightly different regularization method that shrinks weights but doesn’t set them to zero. The weights of features in ridge regression never reached zero, while they often did in

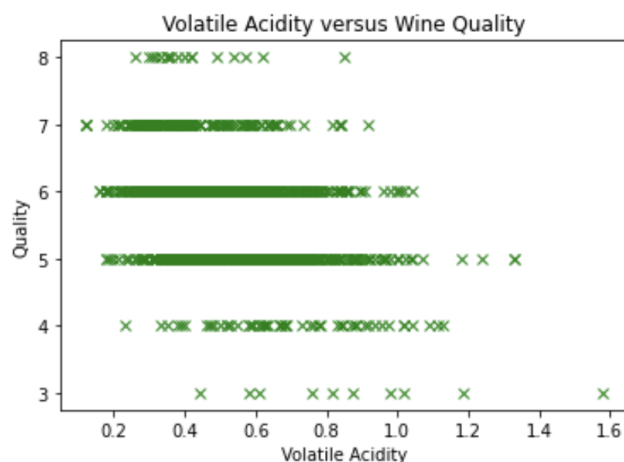


Figure 2: This distribution of wine qualities and their volatile acidity indicates a slight negative correlation between the two variables.

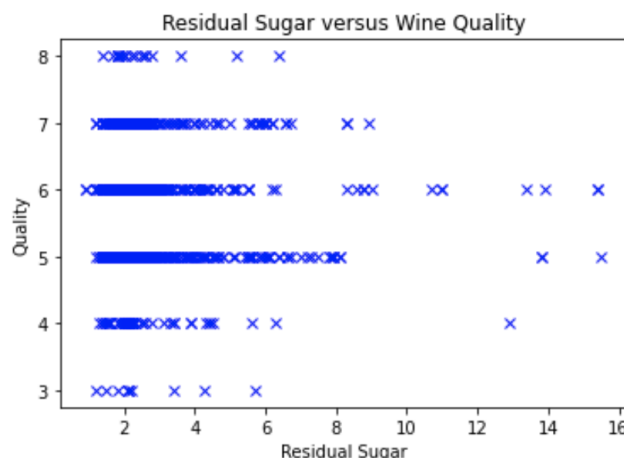


Figure 3: There is no strong correlation between wine quality and residual sugar.

our lasso models, which match the proper functionalities of the two types of regression. Despite converging on different L1 regularization coefficients, the mean squared error of each type of regression was comparable and did not indicate a strong difference in their coefficients of determination. However, we decided to move forward with the lasso approach due to what we know about its built-in feature selector properties.

After settling on lasso regression, we turned to the Lasso structure as a part of *scikit-learn*’s linear modeling package. Lasso regression adopts coordinate descent as its optimizer and tunes learning rate automatically. We performed 5-fold cross validation on our training data, using mean squared error as our scoring metric. In terms of hyperparameters in question, our main focus was on optimizing the L1 regularization coefficient. This would

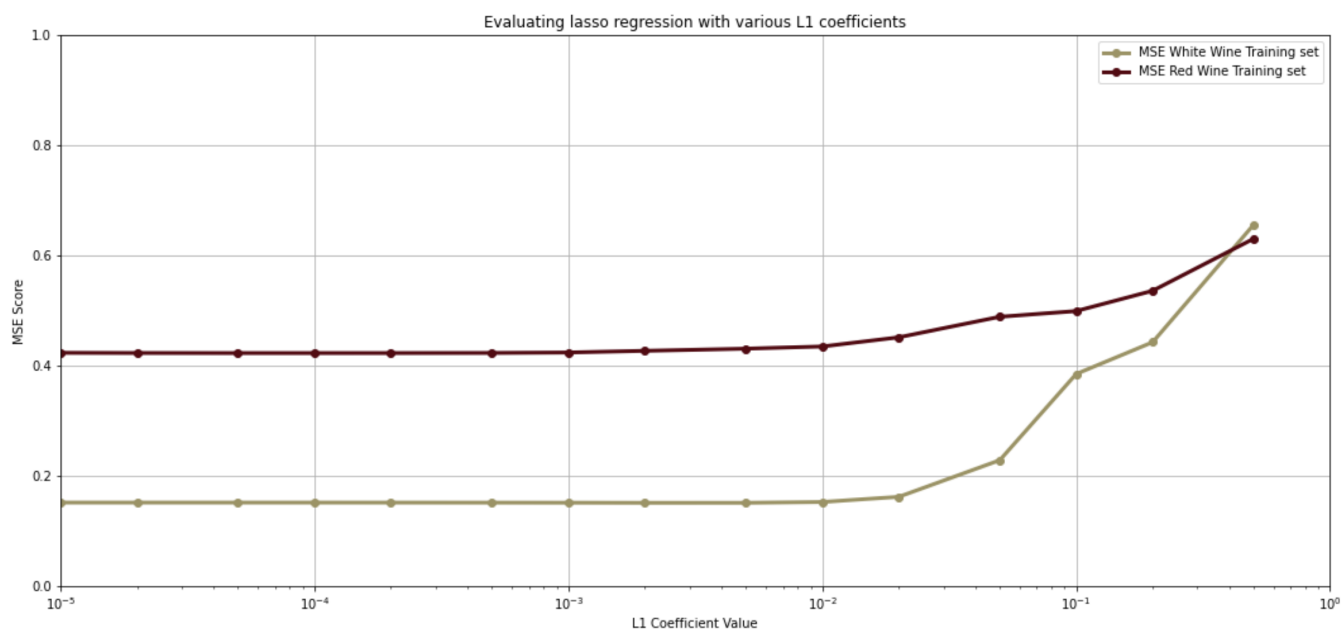


Figure 4: Mean squared error (MSE) by L1 regularization coefficient value for both the red and white wine training sets. The optimal L1 coefficient for the red wine dataset is at 0.00012 with an MSE of 0.424. The optimal L1 coefficient for the white wine dataset is at 0.0024 with an MSE of 0.152. Dark red indicates the red wine dataset, and beige indicates the white wine dataset.

determine the degree at which certain features were rejected from the model. We tuned this value through built-in grid search functionality to optimize the mean squared error of the model (seen in figure 4).

For the red wine dataset, the mean squared error over the 5-fold cross validation data was 0.424 and the optimal L1 coefficient was 0.00012. For the white wine dataset, the mean squared error was 0.152 and the optimal L1 coefficient was 0.0024. After finding the ideal L1 coefficient, we fit our models to the training and test datasets and evaluated them by the coefficient of determination.

4 Results

We fit our two models for red wine and white wine using their respective optimal hyperparameters and found distinct differences between the two types of wine. As for the red wine, the training dataset had a coefficient of determination of 36.56, and the test dataset had a coefficient of determination of 35.06. For the white wine, the training dataset had a coefficient of determination of 27.87, and the test dataset had a coefficient of determination of 28.51. This is tabulated in figure 5. In addition, as seen in figure 4 above, the white wine dataset had a lower mean squared error than the red wine dataset.

This indicates that our model for red wines fit the data

	Red Wine	White Wine
Training Set	36.56	27.87
Test Set	35.06	28.51

Figure 5: Coefficient of determination (R^2) values for both the training and test sets for red and white wine.

slightly better than the model for white wines. One of the biggest differences between the red and white wine datasets (other than the inherent physicochemical differences in the wine) was the size of the dataset. The white wine dataset was roughly three times larger than the red wine dataset, which meant that it trained on far more data. Perhaps we did not address overfitting as dutifully, which may explain the lower mean squared error and coefficient of determination. In terms of the choice of regression model, the minor difference between lasso and ridge regression may indicate the need for further research and experimentation with these models and data.

In addition, the optimized feature coefficients differed between red and white wine data. Values for each coefficient for both datasets is tabulated in figure 6. A lot of the features were weighted similarly between the two types of wine such as residual sugar, free sulfur dioxide, total sulfur dioxide, and density. However, other features like chlorides are drastically different impacts on their respective regression models, indicating chlorides in red wine may be more influential to the overall wine quality.

	Fixed Acidity	Volatile Acidity	Citric Acid	Residual Sugar	Chlorides	Free Sulfur Dioxide
Red Wine	−.016	−0.890	0.227	0.024	−2.176	0
White Wine	−0.023	−1.319	0	0.023	0	0.009

	Total Sulfur Dioxide	Density	pH	Sulphates	Alcohol
Red Wine	−0.003	0	−0.450	0.926	0.261
White Wine	−0.003	0	0	0.496	0.357

Figure 6: Features and their given weights assigned by the regression model for both red and white wine.

5 Broader Impacts

What this model offers the wine industry is a mathematical understanding of the relationships between wine features and rating. People can describe the qualitative features they experience when tasting a wine, but implementing a model such as this provides an opportunity to carry out this type of testing at a faster rate. Additionally, the model that is generated can provide valuable insights into how humans’ rate wine. Individuals could identify through our model which traits affect the rating the greatest. Wineries can profit off this information by adjusting their recipes using this data generated by regressing these data sets.

Despite the automation that machine learning algorithms provide, the implementation and design of these methods depend upon human decision making. Implicit bias that the developer or data source may possess can cause an algorithm to behave improperly, resulting in inaccurate predictions (Mehrabi et al. 2019). By the same token, these biases can worsen inequalities if a machine learning model is rooted in logic that negatively scores specific racial, ethnic, socioeconomic, or other personal traits. Preventing this and other means of biased discrimination must be considered when implementing any machine learning algorithm that has real-world applications (Yapo and Weiss 2018).

For our wine-prediction algorithm, the uses of this tool are specific enough that we are confident there would be minimal unjust consequences that result from our implementation. Because the quality of wine was rated by humans, there is an additional opportunity for personal bias to affect our model’s predictions. One possible inequity could result from our algorithm favoring wines from specific brands or companies. If we were to implement this algorithm in a context that influenced global wine consumption, we would have to perform additional testing to see whether this algorithm might harm the wine market by reducing competition between wineries.

6 Conclusions

Our model sought to predict the rating of wine given eleven quantitative features. While our model accuracy varied depending on the type of wine being rated and the regression model we used, overall it struggled to forecast the new ratings in our testing data. For a model to recreate the qualitative judgements that determined the rating scores, we

hypothesized that the eleven features would assist the model in making a more informed decision. Instead, we found that our lasso regression performed best specifically because it minimized the effect of extra, uncorrelated features.

To continue this work, we recommend introducing a controlled amount of random variability into the coefficients of the best fitting regression model, then plotting the mean squared error of each model against this controlled variability. This plot should show a positive relationship between the variability of the model’s coefficients, and the mean squared error of each model. This relationship would imply that models with coefficients different from ours are less accurate. If this variability and model mean squared error is not positively correlated, our regression model might not have reached the correct relative minimum during descent to accurately model this data.

7 Contributions

E.E. wrote the abstract, introduction, experiments, and results sections. C.S. wrote the data preparation, broader impacts, and conclusion sections. E.E. wrote code that performed data preprocessing steps like train-test split, compared basic lasso and ridge regression models, and created LaTeX tables. C.S. developed and analyzed the visualizations of model performance and hyperparameter tuning, and implemented standardized feature scaling. Both authors helped with hyperparameter tuning and generic model development. Both authors proofread the entire document and all code written.

References

- Buitinck, L.; Louppe, G.; Blondel, M.; Pedregosa, F.; Mueller, A.; Grisel, O.; Niculae, V.; Prettenhofer, P.; Gramfort, A.; Grobler, J.; Layton, R.; VanderPlas, J.; Joly, A.; Holt, B.; and Varoquaux, G. 2013. API design for machine learning software: experiences from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, 108–122.
- Mehrabi, N.; Morstatter, F.; Saxena, N.; Lerman, K.; and Galstyan, A. 2019. A survey on bias and fairness in machine learning.
- Yapo, A., and Weiss, J. 2018. Ethical implications of bias in machine learning.